

Appendix: Share Your Secrets for Privacy! Confidential Forecasting with Vertical Federated Learning

A Datasets and Pre-processing

We make use of the following public datasets in our work: Air Quality [7], SML 2010 [6], PV Power [4], Airline Passengers [5], and Rossman Sales [1]. We do not provide the details on the industrial dataset due to confidentiality agreements. Here we provide a brief description of each dataset and the pre-processing steps applied to the dataset itself. These steps are unconnected with the series pre-processing steps which are part of the STV framework. For all datasets, we scale all features and output values between the range 0 to 1 using a MinMax scaler to ensure that all datasets have the same range of values for comparing the MSE losses.

A.1 Air Quality

The Air Quality dataset contains approximately 9300 samples of multivariate time-series data, with 15 attributes. Five of these are true output values on five gases: Carbon Monoxide (CO), Non Metanic Hydrocarbons (NMHC), Benzene (C6H6), Total Nitrogen Oxides (NOx), and Nitrogen Dioxide (NO2). Exogenous features such as the temperature, ozone levels, and humidity are provided, along with strongly correlated sensor data for each of the five gases. The dataset contains missing values and duplicates, and contains hourly data for each of the five gases.

We preprocess the data by discarding all rows with any missing information and remove duplicate rows. We predict the ground truth values of CO using the other sensor values and information such as temperature and humidity as the exogenous regressors.

A.2 SML 2010

The SML 2010 dataset contains information from a monitor system in a domestic house. It contains approximately 4100 samples with 24 attributes in total, corresponding to 40 days of monitoring data. The attributes contain values such as the indoor and outdoor temperature, lighting levels, Carbon Dioxide levels, relative humidity, rain, wind-speed, etc. We predict the indoor habitation temperature using the others as exogenous features.

A.3 Airline Passengers

Airline Passengers is a small dataset of 145 samples containing the number of international airline passengers (in thousands) on a monthly basis. The exogenous features are also just two: the year

and the month. We predict the number of passengers using the year and month as exogenous regressors.

A.4 PV Power

The PV Power dataset contains around 3100 samples of solar power generation data from each of two power plants over a 34-day period. Attributes include features such as the DC power, AC power, yield, ambient temperature, irradiation levels, and the date and time.

We drop identifiers, empty, and duplicate data. We also drop the DC power attribute, and total yield as these features are very strongly correlated with the AC output. As outputs, we predict the AC power generation using the remaining features as exogenous regressors.

A.5 Rossman Sales

The Rossman Sales dataset contains sales data for 1115 store outlets. The attributes consist of features such as holidays, store type, competitor distance, number of customers, and promotional details among others. We predict the sales of the store with ID 1, using the other features as exogenous regressors.

B Experimental Evaluation

As mentioned in the main text, we use a variation of prequential window testing [2], whereby the entire data is broken into windows of a defined length, each one internally split in an 80-20 train-test ratio. This is illustrated in Figure 1.




Figure 1. Prequential window evaluation with re-training in every window

For each window, we train on the portion allotted for training and forecast the remaining. Within a given training window, we first generate the polynomial by processing the time series as in Algorithm 1 from the main text. Identifying the parameters for generating the

polynomial can be automated using implementations such as `auto.arima`¹.

For each window size, such as 50, 100, 200, 400, we average the MSE loss between the forecasts and the true values across all windows. The average MSE per-window size is given in Table 1, which is an expanded version of Table 1 from the main text.

C SMPC-based XGBoost with VFL

XGBoost [3] is a tree-based gradient-boosting algorithm, that iteratively generates an ensemble of trees by greedily learning a new tree at every step to improve on the earlier one. Each tree has weights assigned to its leaf nodes. When making a prediction for a sample, the weights corresponding to the leaf to which the sample was assigned to are summed to give the final prediction score.

To generate a tree at every iteration, it uses first and second order gradients of the latest predictions, i.e., from the previous tree, in order to set the optimal weights for the new one.

For each sample with index i , the first and second order gradients are denoted as g_i and h_i , respectively. The sum of the gradients of all instances on a particular node is used to set the new weights for it. For example, for node j and corresponding instance set I_j , the accumulated values of g and h are computed as follows: $G_j = \sum_{i \in I_j} g_i$, and $H_j = \sum_{i \in I_j} h_i$. Based on this, for a tree with T nodes, the weights and objective are calculated as follows:

$$w_j = -G_j / (H_j + \lambda) \quad (1)$$

$$obj = -0.5 \times \sum_{j=1}^T ((G_j)^2 / (H_j + \lambda)) + \gamma T \quad (2)$$

, where γ, λ are regularizers. While Equation 1 sets the new weights, Equation 2 is used to identify how to split nodes at each iteration.

With this in mind, using the secret sharing primitives it is possible to compute these functions to extend XGBoost to VFL, which we show in Algorithm 6.

When XGBoost is trained for VFL using secret sharing, each client obtains a local tree with their own weights as shown in Figure 2. In the figure, the weights for clients 1 and 2 are distributed such that their local weights are shares of the weights of the centralized version, i.e., $w_i = w_{i1} + w_{i2} \forall i \in [1, 4]$

The indicator vector s on line 1 of algorithm 1, is a binary vector that is used to point out the location of instances on nodes. We explain this with the help of the example in Figure 3. To calculate the updated weight of the node with instances that have an age greater than 30 (bottom left), we need to find the sum of $\sum_{i \in I_j} g_i$, where $I_j = \{2, 4\}$ (Equation 1). The indicator vector in this case would be $s = [0, 1, 0, 1]$, meaning that nodes 2 and 4 are part of node j . If we have a vector of the gradients, $g = [g_1, g_2, g_3, g_4]$, we can compute $g_2 + g_4$ as $s \odot g$, i.e., the inner product. Under VFL, both the gradients g , and the indicator vector s , exist as secret shares across clients, i.e., $s = \sum_k \langle s_k \rangle^k$, and $g = \sum_k \langle g_k \rangle^k$. Therefore, to compute the product, we can do it using secret shared primitives for matrix multiplication.

The process of split-finding and setting weights is done using the **SecureBuild** function, which makes use of secret sharing primitives to compute the functions in Equation 2 and Equation 1. We defer readers to the implementation in Xie et al. [8] for additional details on this.

¹ https://alkaline-ml.com/pmdarima/modules/generated/pmdarima.arima.auto_arima.html

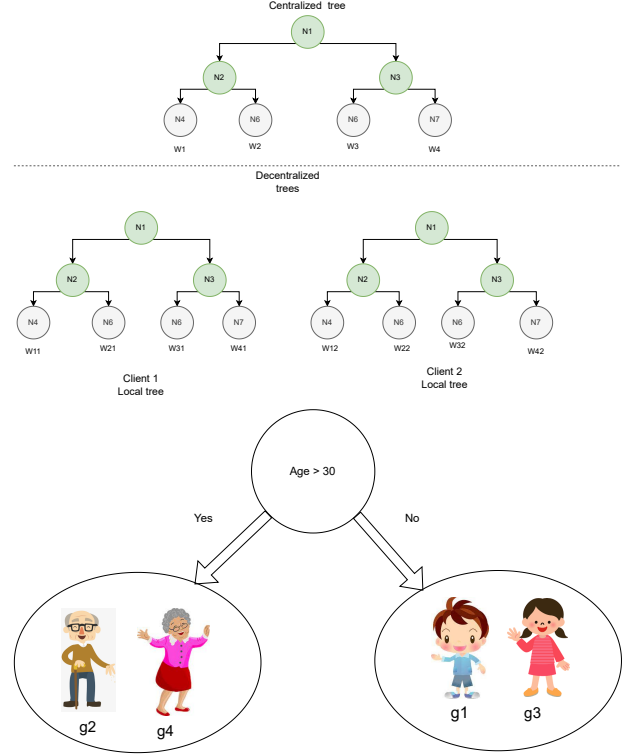


Figure 3. Mapped instances to a node in XGBoost

References

- [1] Rossman store sales, 2015. URL <https://www.kaggle.com/c/rossmann-store-sales>.
- [2] V. Cerqueira, L. Torgo, and I. Mozetič. Evaluating time series forecasting models: An empirical study on performance estimation methods. *Machine Learning*, 109:1997–2028, 2020.
- [3] T. Chen, T. He, M. Benesty, V. Khotilovich, Y. Tang, H. Cho, K. Chen, R. Mitchell, I. Cano, T. Zhou, et al. Xgboost: extreme gradient boosting. *R package version 0.4-2*, 1(4):1–4, 2015.
- [4] A. Kannal. Solar power generation, 2020. URL <https://www.kaggle.com/datasets/anikannal/solar-power-generation-data>.
- [5] C. Kothari. Airline passengers, 2018. URL <https://www.kaggle.com/datasets/chirag19/air-passengers>.
- [6] P. Romeu-Guallart and F. Zamora-Martinez. SML2010. UCI Machine Learning Repository, 2014. DOI: <https://doi.org/10.24432/C5RS3S>.
- [7] S. Vito. Air Quality. UCI Machine Learning Repository, 2016. DOI: <https://doi.org/10.24432/C59K5F>.
- [8] L. Xie, J. Liu, S. Lu, T.-H. Chang, and Q. Shi. An efficient learning framework for federated xgboost using secret sharing and distributed optimization. *ACM Transactions on Intelligent Systems and Technology (TIST)*, 13(5):1–28, 2022.

Algorithm 1 Fit XGBoost [8]

Data: Secretly shared matrices: $\langle X \rangle$, $\langle Y \rangle$, $\langle \hat{Y} \rangle$ across K clients, C_1, C_2, \dots, C_K

Output: Learned tree for the current iteration, one on each client.

- 1: Initialize indicator vector $s \leftarrow 1$ on all C_k
- 2: **if** active client C_1 **then**
- 3: Compute derivatives g, h
- 4: Generate shares $\langle g \rangle$, $\langle h \rangle$, $\langle s \rangle$
- 5: **end if**
- 6: $Tree_k = \text{SecureBuild}(\langle g \rangle^k, \langle h \rangle^k, \langle s \rangle^k)$ on each C_k
- 7: **return** $Tree_k$ on C_k

Dataset	Window size	STV_L	SARIMAX MLE	LSTM	STV_T	$SSSD^{S4}$
Air quality	50	0.00071	0.00110	0.00244	0.00198	0.00270
	100	0.00059	0.00048	0.00055	0.00075	0.00086
	200	0.00066	0.00124	0.00069	0.00067	0.00100
	400	0.00080	0.00068	0.00087	0.00061	0.00144
	Avg.	0.00069	0.00088	0.00114	0.00100	0.00150
SML 2010	50	0.00645	0.00621	0.00343	0.00755	0.01373
	100	0.01958	0.02819	0.02778	0.04305	0.02384
	200	0.00500	0.00662	0.02781	0.01709	0.00336
	400	0.00045	0.00030	0.00960	0.00729	0.00249
	Avg.	0.00787	0.01033	0.01716	0.01875	0.01085
Rossman Sales	50	0.00070	0.00079	0.00464	0.00183	0.00589
	100	0.00064	0.00071	0.00340	0.00496	0.00281
	200	0.00068	0.00086	0.00674	0.00088	0.00233
	400	0.00095	0.00072	0.01079	0.00206	0.00220
	Avg.	0.00074	0.00077	0.00639	0.00243	0.00331
PV Power	25	0.00133	0.00093	0.00950	0.00660	0.00131
	50	0.00360	0.00307	0.01037	0.00258	0.00143
	100	0.00004	0.00063	0.01115	0.00010	0.00267
	200	0.00054	0.00175	0.54227	0.00068	0.00128
	Avg.	0.00138	0.00159	0.14333	0.00249	0.00167
Airline Passengers	60	0.00261	0.00113	0.11651	0.00673	0.00110
	80	0.00450	0.00444	0.00339	0.00628	0.00403
	100	0.00308	0.00066	0.01982	0.00270	0.00798
	120	0.00316	0.00136	0.05164	0.01134	0.00356
	140	0.00185	0.00351	0.01512	0.01332	0.00293
	Avg.	0.00304	0.00222	0.04130	0.00808	0.00392

Table 1. Average normalized MSE values for different public datasets, with different prequential window sizes